

{targets}

What is `{targets}`?



“a Make-like pipeline tool for statistics and data science in R”

- manage a sequence of computational steps
- only update what needs updating
- ensure that the results at the end of the pipeline are still valid

Script-based workflow

01-data.R

```
1 library(tidyverse)
2 data <- read_csv("data.csv", col_types = cols()) %>%
3   filter(!is.na(Ozone))
4 write_rds(data, "data.rds")
```

02-model.R

```
1 library(tidyverse)
2 data <- read_rds("data.rds")
3 model <- lm(Ozone ~ Temp, data) %>%
4   coefficients()
5 write_rds(model, "model.rds")
```

03-plot.R

```
1 library(tidyverse)
2 model <- read_rds("model.rds")
3 data <- read_rds("data.rds")
4 ggplot(data) +
5   geom_point(aes(x = Temp, y = Ozone)) +
6   geom_abline(intercept = model[1], slope = model[2])
7 ggsave("plot.png", plot)
```

Problems with script-based workflow

- **Reproducibility:** if you change something in one script, you have to remember to re-run the scripts that depend on it
- **Efficiency:** that means you'll usually rerun all the scripts even if they don't depend on the change
- **Scalability:** if you have a lot of scripts, it's hard to keep track of which ones depend on which
- **File management:** you have to keep track of which files are inputs and which are outputs and where they're saved

{targets}: The basics

{targets} workflow

R/functions.R

```
1 get_data <- function(file) {  
2   read_csv(file, col_types = cols()) %>%  
3   filter(!is.na(Ozone))  
4 }  
5  
6 fit_model <- function(data) {  
7   lm(Ozone ~ Temp, data) %>%  
8   coefficients()  
9 }  
10  
11 plot_model <- function(model, data) {  
12   ggplot(data) +  
13     geom_point(aes(x = Temp, y = Ozone)) +  
14     geom_abline(intercept = model[1,1], slope = model[2,1])
```

{targets} workflow

_targets.R

```
1 library(targets)
2
3 tar_source()
4 tar_option_set(packages = c("tidyverse"))
5
6 list(
7   tar_target(file, "data.csv", format = "file"),
8   tar_target(data, get_data(file)),
9   tar_target(model, fit_model(data)),
10  tar_target(plot, plot_model(model, data))
11 )
```

Run `tar_make()` to run pipeline



`use_targets()` will generate a `_targets.R` script for you to fill in.

{targets} workflow

Targets are “hidden” away where you don’t need to manage them

```
|— _targets.R
|— data.csv
|— R/
|   |— functions.R
|— _targets/
|   |— objects
|       |— data
|       |— model
|       |— plot
```



You can of course have multiple files in `R/`; `tar_source()` will source them all

My typical workflow with `{targets}`

1. Read in some data and do some cleaning until it's in the form I want to work with.
2. Wrap that in a function and save the file in `R/`.
3. Run `use_targets()` and edit `_targets.R` accordingly, so that I list the data file as a target and `clean_data` as the output of the cleaning function.
4. Run `tar_make()`.
5. Run `tar_load(clean_data)` so that I can work on the next step of my workflow.
6. Add the next function and corresponding target when I've solidified that step.



I usually include `library(targets)` in my project `.Rprofile` so that I can always call `tar_load()` on the fly

_targets.R tips and tricks

```
1 list(  
2   tar_target(  
3     data_file,  
4     "data/raw_data.csv",  
5     format = "file"  
6   ),  
7   tar_target(  
8     raw_data,  
9     read.csv(data_file)  
10  ),  
11  tar_target(  
12    clean_data,  
13    clean_data_function(raw_data)  
14  ),  
15  )
```



_targets.R tips and tricks

```
1 preparation <- list(  
2   ...,  
3   tar_target(  
4     clean_data,  
5     clean_data_function(raw_data)  
6   )  
7 )  
8 modeling <- list(  
9   tar_target(  
10    linear_model,  
11    linear_model_function(clean_data)  
12  ),  
13  ...  
14 )
```



By grouping the targets into lists, I can easily comment out chunks of the pipeline to not run the whole thing

_targets.R tips and tricks

```
1 ## prepare ----
2 prepare <- list(
3   ### cleanData.csv ----
4   tar_target(
5     cleanData.csv,
6     file.path(path_to_data,
7               "cleanData.csv"),
8     format = "file"
9   ),
10  ### newdat ----
11  tar_target(
12    newdat,
13    read_csv(cleanData.csv,
14             guess_max = 20000)
```

prepare	raw
cleanData.csv	log_res_raw
newdat	mult_res_raw
fulldat	ctc_res_raw
flow	tte_res_raw
loglinear	tte_res_raw_very
log_dat	all_res_raw
log_dat_covid_mild...	imputation
log_dat_covid	mice_params
log_est	cols_to_impute
log_est_severity	fulldat_imputed
loglinear_spon	log_dat_imputed
log_est_spon	mice_comparison
log_est_severity_spon	tabs_figs
log_est_ind	descriptive_table
log_est_severity_ind	descriptive_table_o...
mult_est	outcomes_table
mult_est_severity	risk_plot
ctc	tte_plot
ctc_dat	followup_tab
cases	testing_plot
controls	all targets



In big projects, I comment my `_targets.R` file so that I can use the RStudio outline pane to navigate the pipeline (`my buggy function`)

Key `{targets}` functions

- `use_targets()` gets you started with a `_targets.R` script to fill in
- `tar_make()` runs the pipeline and saves the results in `_targets/objects/`
- `tar_make_future()` runs the pipeline in parallel¹
- `tar_load()` loads the results of a target into the global environment (e.g., `tar_load(clean_data)`)
- `tar_read()` reads the results of a target into the global environment (e.g., `dat <- tar_read(clean_data)`)
- `tar_visnetwork()` creates a network diagram of the pipeline
- `tar_outdated()` checks which targets need to be updated
- `tar_prune()` deletes targets that are no longer in `_targets.R`
- `tar_destroy()` deletes the `.targets/` directory if you need to burn everything down and start again

¹ Note: `{targets}` is moving to a new distributed computing strategy using `foreach`

Advanced {targets}

“target factories”



repo status **Active** tarchetypes makes it easy to add certain kinds of common tasks to reproducible pipelines. Most of its functions create families of targets for parameterized R Markdown, simulation studies, and other general-purpose scenarios.



repo status **Active** stantargets is a workflow framework for Bayesian data analysis with cmdstanr. With concise, easy-to-use syntax, it defines versatile families of targets tailored to Bayesian statistics, from a single MCMC run with postprocessing to large simulation studies.



repo status **Active** Like stantargets, jagstargets is a workflow framework for Bayesian data analysis, with support for both single MCMC runs and large-scale simulation studies. It invokes JAGS through the R2jags package, which has nice features such as the ability to parallelize chains across local R processes.

`{tarchetypes}`: reports

Render documents that depend on targets loaded with `tar_load()` or `tar_read()`.

- `tar_render()` renders an R Markdown document
- `tar_quarto()` renders a Quarto document (or project)



It can't detect dependencies like `tar_load(ends_with("plot"))`

What does `report.qmd` look like?

```
1 ---
2 title: "My report"
3 ---
4 ```{r}
5 library(targets)
6 tar_load(results)
7 tar_load(plots)
8 ```
9 There were `r results$n` observations with a mean age of `r result
10 ```{r}
11 library(ggplot2)
12 plots$age_plot
13 ```
```

Because `report.qmd` depends on `results` and `plots`, it will only be re-rendered if either of those targets change.



The `extra_files =` argument can be used to force it to depend on additional non-target files

{tarchetypes}: branching

Using data from the National Longitudinal Survey of Youth,

`_targets.R`

```
1 library(targets)
2 library(tarchetypes)
3 tar_source()
4
5 targets_setup <- list(
6   tar_target(
7     csv,
8     "data/nlsy.csv",
9     format = "file"
10  ),
11  tar_target(
12    dat,
13    readr::read_csv(csv,
14      show_col_types = FALSE)
15  )
16 )
```

`R/functions.R`

```
1 model_function <- function(outcome_var,
2                             sex_val, dat) {
3
4   lm(as.formula(paste(outcome_var,
5     " ~ age_bir + income + factor(region)")),
6     data = dat,
7     subset = sex == sex_val)
8 }
9
10 coef_function <- function(model) {
11   coef(model)[["age_bir"]]
12 }
```

we want to investigate the relationship between age at first birth and hours of sleep on weekdays and weekends among moms and dads separately

Option 1

Create (and name) a separate target for each combination of sleep variable ("sleep_wkdy", "sleep_wknd") and sex (male: 1, female: 2):

```
1 targets_1 <- list(  
2   tar_target(  
3     model_1,  
4     model_function(outcome_var = "sleep_wkdy", sex_val = 1, dat =  
5   ),  
6   tar_target(  
7     coef_1,  
8     coef_function(model_1)  
9   )  
10 )
```

... and so on...

Option 2

Use `tarchetypes::tar_map()` to map over the combinations for you (static branching):

```
1 targets_2 <- tar_map(  
2   values = tidyr::crossing(  
3     outcome = c("sleep_wkdy", "sleep_wknd"),  
4     sex = 1:2  
5   ),  
6   tar_target(  
7     model_2,  
8     model_function(outcome_var = outcome, sex_val = sex, dat = dat  
9   ),  
10  tar_target(  
11    coef_2,  
12    coef_function(model_2)  
13  )  
14 )
```

Option 2, cont.

Use `tarchetypes::tar_combine()` to combine the results of a call to `tar_map()`:

```
1 combined <- tar_combine(  
2   combined_coefs_2,  
3   targets_2[["coef_2"]],  
4   command = vctrs::vec_c(!!!.x),  
5 )  
6 tar_read(combined_coefs_2)
```

`command = vctrs::vec_c(!!!.x)` is the default, but you can supply your own function to combine the results

Option 3

Use the `pattern =` argument of `tar_target()` (dynamic branching):

```
1 targets_3 <- list(  
2   tar_target(  
3     outcome_target,  
4     c("sleep_wkdy", "sleep_wknd")  
5   ),  
6   tar_target(  
7     sex_target,  
8     1:2  
9   ),  
10  tar_target(  
11    model_3,  
12    model_function(outcome_var = outcome_target, sex_val = sex_target,  
13    pattern = cross(outcome_target, sex_target)  
14  ),  
15)
```

Branching

Dynamic

Pipeline creates new targets at runtime.

Cryptic target names.

Scales to hundreds of branches.

No metaprogramming required.

Static

All targets defined in advance.

Friendly target names.

Does not scale as easily for `tar_visnetwork()` etc.

Familiarity with metaprogramming is helpful.

Branching

- The book also has an example of using metaprogramming to map over different functions
 - i.e. fit multiple models with the same arguments
- Static and dynamic branching can be combined
 - e.g. `tar_map(values = ..., tar_target(..., pattern = map(...)))`
- Branching can lead to slowdowns in the pipeline (see book for suggestions)

{tarchetypes}: repetition

`tar_rep()` repeats a target multiple times with the same arguments

```
1 targets_4 <- list(  
2   tar_rep(  
3     bootstrap_coefs,  
4     dat |>  
5       dplyr::slice_sample(prop = 1, replace = TRUE) |>  
6       model_function(outcome_var = "sleep_wkdy", sex_val = 1, dat  
7       coef_function()),  
8     batches = 10,  
9     reps = 10  
10  )  
11 )
```

The pipeline gets split into `batches` x `reps` chunks, each with its own random seed

{tarchetypes}: mapping over iterations

```
1 sensitivity_scenarios <- tibble::tibble(  
2   error = c("small", "medium", "large"),  
3   mean = c(1, 2, 3),  
4   sd = c(0.5, 0.75, 1)  
5 )
```

`tar_map_rep()` repeats a target multiple times with different arguments

```
1 targets_5 <- tar_map_rep(  
2   sensitivity_analysis,  
3   dat |>  
4     dplyr::mutate(sleep_wkdy = sleep_wkdy + rnorm(nrow(dat), mean,  
5     model_function(outcome_var = "sleep_wkdy", sex_val = 1, dat =  
6     coef_function() |>  
7     data.frame(coef = _),  
8   values = sensitivity_scenarios,
```

```
9     batches = 10,  
10     reps = 10  
11 )
```

{tarchetypes}: mapping over iterations

```
1 tar_read(sensitivity_analysis) |> head()
```

Ideal for sensitivity analyses that require multiple iterations of the same pipeline with different parameters

```
1 tar_read(sensitivity_analysis) |>
2   dplyr::group_by(error) |>
3   dplyr::summarize(q25 = quantile(coef, .25),
4                   median = median(coef),
5                   q75 = quantile(coef, .75))
```

Conclusion

- `{targets}` is a great tool for managing complex workflows
- `{tarchetypes}` makes it even more powerful
- The [user manual](#) is a great resource for learning more

